

Highlights

Explainable production-adjusted electrical-energy baselines for industrial feed pelleting

Vladimir Sousa Santos^{}, Juan José Cabello Eras^{}, Eder Elías Herrera Rodiño, Francisco Manuel Arrabal-Campos^{}

- Production-adjusted electricity baseline from 2,745 real feed-pelleting records
- Frozen Ridge model cuts test MAE by 14.42% relative to line-specific SEC
- TFT, N-HiTS, TimesFM 2.5 and Chronos-2 add no accuracy on identical plant records
- Exact additive decomposition traces each kWh estimate to recorded production inputs
- Fixed 90%/95% residual intervals cover 85.96% and 93.23% of later test records

Explainable production-adjusted electrical-energy baselines for industrial feed pelleting

Vladimir Sousa Santos^a, Juan José Cabello Eras^b, Eder Elías Herrera Rodiño^c,
Francisco Manuel Arrabal-Campos^{d,*}

^a*Department of Energy, Universidad de la Costa, Calle 58 # 55–66, Barranquilla, Colombia*

^b*Department of Mechanical Engineering, Faculty of Engineering, Universidad de
Córdoba, Carrera 6 No. 77–305, Montería, Colombia*

^c*Master's Programme in Mechanical Engineering, Universidad de Córdoba, Carrera 6 No.
77–305, Montería, Colombia*

^d*Department of Engineering, Research Centre CIAIMBITAL, Escuela Superior de Ingeniería,
University of Almería, Carretera de Sacramento s/n, La Cañada de San
Urbano, 04120, Almería, Spain*

Abstract

This study asks whether an explainable, production-adjusted baseline can state how much electricity a feed-pelleting operation should have used, and whether operating histories, modern temporal models or pretrained foundation models add anything to it. Electricity use in pelleting varies with the quantity produced, the product, the line and the batch organization, so neither raw consumption nor a single specific energy consumption (SEC) figure says much about performance. We use 2,745 real records from three industrial lines (2,742 with positive energy) with dosed mass, dosing-batch count, product characteristics, line identity and calendar variables as inputs. Seven methods were compared on 800 chronologically ordered development targets; the selected Ridge regression was frozen, with its interval calibration, and evaluated once against line-specific SEC on 798 later records. Eight current-input, history-enabled, temporal and pretrained configurations were compared on 584 matched targets with 16 completed same-line operations as context. Ridge reduced test MAE from 61.29 to 52.45 kWh, 14.42% below SEC,

*Corresponding author.

Email addresses: vsousa1@cuc.edu.co (Vladimir Sousa Santos^a),
juancabelloe@correo.unicordoba.edu.co (Juan José Cabello Eras^b),
eeherrera@correo.unicordoba.edu.co (Eder Elías Herrera Rodiño), fmarrabal@ual.es
(Francisco Manuel Arrabal-Campos^d)

on all three lines; this is lower prediction error, not an energy saving. The Temporal Fusion Transformer reached 53.65 kWh against 53.55 kWh for current-input Ridge, with no evidence of paired improvement, and N-HITS, TimesFM 2.5 and Chronos-2 were less accurate. Every estimate decomposes exactly into its inputs, and fixed 90% and 95% residual intervals covered 85.96% and 93.23% of the test records, with one line systematically underestimated. The baseline is ready for retrospective assessment and for a digital twin; prospective use needs verified inputs, line-specific acceptance criteria and a new consecutive period.

Keywords: feed pelleting, industrial energy baseline, production adjustment, chronological validation, prediction intervals, explainable Ridge regression

Nomenclature

Symbols

y_i, E_i	pelleting electrical energy of record i (kWh)
$\hat{y}_i$	point estimate of y_i (kWh)
m_i	dosed mass of record i (t)
b_i	number of dosing batches of record i
ℓ_i, c_i, p_i, s_i	pelleting line, product subfamily, presentation and bag/bulk category of record i
h_i, d_i	origin hour (h) and weekday (0–6) of record i
t_i	analytical origin of record i : earliest recorded stage start
a_j	recorded pelleting completion of record j
Δ_i	recorded pelleting duration of record i (h)
P_i^{avg}	duration-based average-power proxy of record i (kW)
z_i	standardized and one-hot encoded input vector of record i
$\hat{b}, \hat{\beta}$	fitted intercept and coefficient vector of the Ridge model
$\hat{s}_\ell$	training-period specific energy consumption of line ℓ (kWh/t)
$\mathcal{T}, \mathcal{P}$	training set and partition interval
$\mathcal{H}_i^{(16)}, \mathcal{H}_i$	16-record operating history and recent-residual history of target i
C_{ig}	contribution of variable group g to the estimate of record i (kWh)
δ_i	shrunk recent-residual correction (kWh)
q_α	calibrated half-width at miscoverage level α (kWh)
L_i, U_i	lower and upper interval limits of record i (kWh)
r_i^{excess}, B_i	signed excess over the estimate (kWh) and out-of-interval flag

Abbreviations

EnB, EnPI	energy baseline, energy performance indicator
IPMVP	International Performance Measurement and Verification Protocol
MAE, RMSE, WAPE	mean absolute error, root mean squared error, weighted absolute percentage error
IS	interval score
MLP	multilayer perceptron
SEC	specific energy consumption (kWh/t)
TFT	Temporal Fusion Transformer
N-HITS	Neural Hierarchical Interpolation for Time Series
XReg	covariate-regression adapter of TimesFM
H1, H2, H3	post-hoc extensions: line-dependent mass effect, product-dependent mass effect, recent-residual correction
GPU, CPU	graphics and central processing unit

1. Introduction

Two pelleting operations can use very different amounts of electricity for entirely ordinary reasons: one processes more material, runs a different product, or takes place on another line. When a plant asks whether an operation used more energy than it should have, those differences must be accounted for before anything can be said about efficiency. The usual answer in feed manufacturing is a specific energy consumption (SEC) in kWh/t, which is simple and interpretable but cannot separate the effects of product mix, line assignment, processed mass and batch structure. Energy management standards make the alternative explicit: ISO 50001 asks for energy performance to be judged against a baseline normalized for the relevant variables, and ISO 50006 and the IPMVP describe how such baselines and performance indicators are built and used [1, 2, 3]. Electricity is also a large share of the sector’s energy bill: food processing in the United States consumes about 75 TWh of electricity, and electricity accounts for roughly 30% of the reported energy use of UK food and beverage manufacturing [4, 5, 6].

Pelleting is a good place to build such a baseline, because its demand is strongly conditioned by what is processed. Pellet quality, conditioning, feed composition, steam application, temperature, particle characteristics and operating conditions all shape pelleting behaviour and mechanical energy demand [7, 8, 9, 10, 11]. Yet those studies aim at material properties, process behaviour or pellet quality; none builds and tests, chronologically, an operation-level electricity baseline from real plant records. That is the gap this paper addresses, and the next subsection places it against the closest published work.

1.1. Related work

Industrial energy baselines and performance indicators. Regression-based baselines for measurement and verification go back at least to the monthly-data models of Kissock and Eger [12], and automated baselines on interval data showed that predictive accuracy and honest uncertainty decide whether a saving can be claimed [13, 14]. Key-performance-indicator frameworks translate the same idea to production facilities [15]. In the feed industry itself, Sánchez et al. [16] built motor-level efficiency assessments, energy baselines and performance indicators for a balanced-feed plant (about 10 kWh/t at 500 t/day, with an estimated saving potential of 364,000 kWh/year), but at an aggregated process level rather than per operation. Wigger et al. [17] showed how AI models can be integrated into the energy management system of a manufacturing company from existing data resources, with consumption forecasting and anomaly detection as the use cases; their object is the company’s consumption series, not the individual operation and what it processed.

Data-driven models of process energy. Regularized regression gives stable, interpretable estimates when predictors are correlated [18], and tree ensembles and neural networks capture nonlinearities and interactions [19]. Closest to our setting, Ekwaro-Osire et al. [20] compared linear regression, tree-based methods and a multilayer perceptron on about 1,100 production batches with 67 input characteristics from an industrial compound-feed pelleting process and found that more complex models did not necessarily reduce error. That study relied on shuffled cross-validation, so it could not say how the models behave on a later period, and it did not consider operating histories, pretrained temporal models, predictive uncertainty or record-level traceability. Panjapornpon et al. [21] predicted the energy efficiency of a petrochemical process with an explainable transfer-learning LSTM made robust to sensor faults, and Mohammed et al. [22] combined multilinear regression with neural components for short-term load forecasting; both work on regularly sampled signals rather than on heterogeneous completed operations.

Temporal and foundation models. The Temporal Fusion Transformer (TFT) [23] and N-HiTS [24] are the reference neural forecasters with covariates, and foundation models such as TimesFM and Chronos transfer representations learned from large time-series collections to new tasks [25, 26]. Benchmarks in this journal have begun to map where they help: the FETS benchmark evaluates foundation models against task-specific learners across 54 energy datasets [27]; zero-shot Chronos and TimesFM variants have been compared with tuned LSTMs and conventional models

for electricity prices [28, 29]; TimesFM has been assessed for battery degradation across datasets [30]; and decision-focused fine-tuning has been proposed for feeder optimization [31]. Comparative frameworks also stress the evaluation design itself (fixed-origin, expanding and rolling windows against linear and naive baselines [32]; context length and horizon effects [33]; transfer and zero-shot strategies [34]). All of these operate on regularly sampled series. Industrial pelleting records are different: consecutive completed operations have different masses, products, lines, start times and durations, so the prediction origin and the information available at that moment must be fixed before any model can be compared [35], and random or shuffled partitions inflate apparent performance [36, 37, 38].

Explainability and uncertainty. Explainable AI for energy systems has a growing literature [39, 40, 41], from frameworks that validate what data-driven power-curve models have learned [42] and sparse models that recover governing relations from plant data [43], to explainable drift detection in consumption forecasts [44]. Prediction intervals complete the picture: conformal methods have been compared for electricity prices and solar power [45, 46], and their guarantees are known to weaken under temporal dependence and drift [47, 48, 49]; coverage must therefore be read together with width and checked by subgroup [50, 51]. Digital twins, finally, are the natural destination of a validated baseline [52, 53], provided the statistical component is trustworthy.

Table 1 summarizes how the closest studies relate to the present one. No previous work combines real operation-level records, chronological validation with a frozen test, an exactly decomposable baseline, residual-calibrated intervals checked by line and a matched comparison with temporal and foundation models.

1.2. Objective and contributions

The objective of this study is to establish, on real operational records of three industrial feed-pelleting lines, whether an explainable, production-adjusted electricity baseline predicts the energy of a completed operation better than the line-specific SEC reference used in practice, and whether histories of completed operations, modern temporal models or pretrained foundation models add anything to it. Three things are new:

1. **A chronologically validated, production-adjusted baseline for heterogeneous operations.** Seven estimators are compared on 800 chronological development targets, the selected Ridge regression is frozen with its interval calibration, and it is evaluated once on 798 later records against line-specific

Table 1: The closest published studies and what this one adds. “Operation-level” means one record per completed industrial operation with its production descriptors; “frozen test” means a model fixed before a later period is evaluated.

Study	Data and target	Production adjusted	Frozen chronological test	Explainability / uncertainty	Temporal or foundation models
Kissock and Eger [12]	monthly plant energy; weather, production	yes	partly	linear terms / no	no
Granderson et al. [13]; Gallagher et al. [14]	interval meter data; buildings, industrial site	yes	yes	no / savings uncertainty	no
Sánchez et al. [16]	feed plant; motor efficiency, EnPIs	aggregated kWh/t	no	no / no	no
Ekwaro-Osire et al. [20]	$\approx 1,100$ pelleting batches, 67 inputs	yes	no (shuffled CV)	no / no	no
Wigger et al. [17]	company consumption series	no	forecast split	no / no	conventional forecasters
Panjapornpon et al. [21]	petrochemical signals; energy efficiency	process variables	no	post-hoc XAI / fault uncertainty	LSTM
Letzgus and Müller [42]	wind turbine SCADA; power curve	ambient conditions	yes	post-hoc XAI / no	no
FETS [27]; Miseta et al. [28]; Lin and Ti [29]	regular energy series; prices	covariates where available	yes	no / partly	foundation vs. task-specific
O’Connor et al. [45]; Nam et al. [46]	electricity prices; solar power	no	yes	no / conformal, global	conventional, Transformer
This study	2,745 real pelleting operations; kWh per operation	yes	yes	exact additive / by line and week	TFT, N-HITS, TimesFM 2.5, Chronos-2 on identical records

SEC. Previous feed-industry baselines are aggregated [16] or validated by shuffling [20].

2. **A matched test of operating histories, temporal models and foundation models on irregular industrial records.** TFT, N-HITS, TimesFM 2.5 and Chronos-2 receive exactly the same 16 completed same-line operations and the same current descriptors as the linear control, on 584 identical targets, with the information available at the origin fixed in advance. Existing foundation-model benchmarks concern regularly sampled series [27, 28, 29].
3. **Record-level explainability and calibrated uncertainty that survive the chronological test.** Every estimate decomposes exactly into production, product, line and calendar terms, and residual-calibrated 90% and 95% intervals are assessed by line and by week, which is what distinguishes an unusual operation from a product mix.

The validated scope is retrospective assessment at the closure of recorded operations;

prospective monitoring is a later application that requires confirming predictor availability at the intended origin. For a plant, the result is a baseline that separates the ordinary variability of electricity use from deviations worth a look, and that can serve as the statistical component of an energy-management platform or digital twin. Section 2 describes the records, models and protocol, Section 3 the results, Section 4 their implications and limits, and Section 5 concludes.

2. Materials and methods

2.1. Study design

This is a retrospective observational model-development study on recorded operations of three industrial feed-pelleting lines. It estimates the expected electricity of completed operations rather than forecasting a continuously sampled power signal, and its observational unit is an operation represented by one row of the plant database. The primary evaluation develops and selects models on temporally ordered validation periods, freezes the selected model and evaluates it against operational reference baselines in a later test period, and assesses predictive performance and interval coverage overall and by line. Complementary analyses examine short histories of completed same-line operations with locally trained temporal architectures and pretrained time-series models, line- and product-specific mass interactions, a recent-residual correction, rolling updates of the interval widths, and numerical agreement between computational implementations. None of them replaces the frozen model. Figure 1 summarizes the workflow.

2.2. Industrial records and observational unit

The plant workbook holds 2,745 records and 32 named variables from three lines: 730 from Pellet 1, 1,148 from Pellet 2 and 867 from Pellet 3, with recorded stage starts from 2 January to 1 April 2025 and the last pelleting completion on 2 April. The plant information system generated each row as a consolidated record of one production operation: dosed mass, dosing-batch count, product classifications, stage timestamps, pelleting electrical energy (KWH PELET), steam-related variables and operational summaries. There is no verified production-order identifier linking every stage to a single physical batch and no separate meter history from which the energy value could be reconstructed, so a “recorded production operation” is one row of the workbook, and source-row identifiers are preserved throughout. Records with overlapping timestamps were kept when their identifiers differed, because overlap alone is no evidence of duplication; overlaps are characterized by line and treated as a data-quality descriptor.




Figure 1: Methodological workflow. The primary evaluation comprises chronological model development, frozen test evaluation and line-specific assessment; the complementary analyses examine short-context temporal models, line- and product-specific interactions, recent-residual correction, rolling interval updates and computational diagnostics.

2.3. Target, predictors and preprocessing

The response y_i is KWH PELET, the electrical energy attributed to pelleting record i in kWh; it excludes the separately recorded electricity of milling and mixing and the thermal energy supplied as steam. Dosed mass m_i comes from PESO DOS (kg) converted to tonnes. Energy, a duration-based average-power proxy and specific energy consumption are defined as

$$E_i = y_i \text{ [kWh]}, \quad P_i^{\text{avg}} = \frac{y_i}{\Delta_i} \text{ [kW]}, \quad \text{SEC}_i = \frac{y_i}{m_i} \text{ [kWh/t]}, \quad (1)$$

where Δ_i is the recorded pelleting duration in hours. P_i^{avg} is derived, not measured; E_i is the modelling target and SEC is used only to build the reference baselines.

The analytical origin t_i is the earliest recorded start among milling, mixing and pelleting; it assigns records to partitions and decides whether a historical outcome was completed before a target. The predictors are dosed mass, dosing-batch count, pelleting line, product subfamily, presentation, bag/bulk category, origin hour and origin weekday (Table 2). With b_i the number of dosing batches, h_i the origin hour

and $d_i \in \{0, \dots, 6\}$ the weekday from Monday, the numerical input vector is

$$x_i^{\text{num}} = \left(m_i, b_i, \sin \frac{2\pi h_i}{24}, \cos \frac{2\pi h_i}{24}, \sin \frac{2\pi d_i}{7}, \cos \frac{2\pi d_i}{7} \right). \quad (2)$$

The six numerical coordinates are standardized with training-data parameters only and the four categorical variables are one-hot encoded with training vocabularies, giving 22 columns in the Ridge design matrix; a level absent from a training period produces no active indicator. Grouped Ridge contributions therefore depend partly on this coding and are read as components of the fitted representation, not as invariant physical effects. The analysis is retrospective and conditional on the consolidated production characteristics of each operation: current-record durations, stage energies and the steam, temperature and current summaries are excluded as inputs because their availability at the origin cannot be established, and recorded pelleting completion is taken as the availability time of historical outcomes. A prospective implementation would have to restrict the model to descriptors verified as available at the intended origin.

2.4. Data integrity and temporal consistency

All 32 source fields were reconciled with the analytical table across the 2,745 records: the 87,840 cell-level comparisons satisfied the predefined numerical and whitespace tolerances. Three records (0.11%) with $\text{KWH PELET}=0$ were excluded from supervised modelling, without imputation. Table 3 lists the main checks; the counts describe different characteristics of possibly the same records. Median pelleting durations are about 145, 96 and 123 minutes on Pellet 1, 2 and 3; milling completion falls after pelleting start in 528 records and mixing completion in 376. These patterns are characteristics of consolidated plant records rather than exclusion criteria. The workbook contains event-level summaries, not dense power trajectories or verified active/inactive labels; a separately supplied set of synthetic ten-second series was kept apart from the industrial evidence.

2.5. Chronological evaluation design

Records are assigned by origin to development (before 22 February 2025), calibration (22 February to 5 March) and testing (from 6 March); the last operation completes on 2 April (Figure 2). Development and calibration eligibility require a positive outcome, a valid duration, complete inputs and completion before the next boundary, which retains 1,542 of 1,556 development candidates and 380 of 391 calibration candidates (the 25 exclusions of Table 3); all 798 test candidates are retained, including the five operations longer than 24 hours.

Table 2: Model inputs: source fields, preprocessing, recorded support and conditions for prospective use. Original Spanish field names are retained for traceability; counts and ranges refer to all 2,745 source records.

Input / units	Source field and preprocessing	Interpretation and recorded support	Prospective use
Dosed mass m_i / t	PESO DOS (kg); divided by 1,000	Material entering the operation; 2.69–36.78 t, median 14.98 t; the mass denominator used throughout	Possible if the mass is available at the intended origin
Dosing batches b_i / count	N_BACHES_DOS; numerical	One to ten, median five; describes how the mass is aggregated and is correlated with it; not a physical-cycle identifier	Possible when the batch plan or completed count is available
Pelleting line ℓ_i / category	PELLT; one-hot	Pellet 1, 2 and 3: 730, 1,148 and 867 records; the separate field LINEA holds product families	Possible when the line assignment is confirmed
Product subfamily c_i / category	SUBLINEA; one-hot	Nine labels spanning poultry, swine, cattle, equine and rabbit products; a classification, not a measured formulation	Possible when the product classification is known
Presentation p_i / category	PRESENTACION; one-hot	PELETIZADO (pelleted): 2,065; CROMBELIZADO (crumbled): 680; may also reflect product allocation	Possible when the presentation is known
Bag/bulk s_i / category	DES_PRODUCTO_ TIPO; one-hot	ENSAQUE (bagged): 2,420; GRANEL (bulk): 325; the analytical alias salida_programada identifies output form, not a departure time	Possible when the output form is known
Origin hour / dimensionless pair	Earliest recorded stage start; sine and cosine	Two coordinates avoid a discontinuity between 23:59 and 00:00; seconds are not used	Available at or immediately after the origin
Origin weekday / dimensionless pair	Same origin; sine and cosine	Two coordinates represent recurring weekly associations, not operational state	Available at the origin

Table 3: Data-integrity and temporal-consistency checks.

Check	Count	Analytical treatment
Source records	2,745	Reconciled with the analytical table.
Records with zero pelleting energy	3	Excluded from supervised targets; not imputed.
Imputed auxiliary values	15	Retained in the source copy; affected variables were not model inputs.
Within-line interval overlaps	351	Retained as distinct source records and characterized by pelleting line.
Pellet 1 overlaps	75	Included in the line-level overlap count.
Pellet 2 overlaps	131	Included in the line-level overlap count.
Pellet 3 overlaps	145	Included in the line-level overlap count.
Test records exceeding 24 hours	5	Retained in the primary test population.
Pre-test eligibility exclusions	25	Fourteen development and eleven calibration candidates.

Model development uses four expanding training blocks with validation weeks beginning on 25 January, 1 February, 8 February and 15 February 2025: training sizes 705, 924, 1,121 and 1,334, validation sizes 208, 187, 205 and 200, i.e. 800 common targets over 28 origin dates. Training outcomes must be completed before each validation period starts, and validation outcomes before the week ends, so that the expanding blocks remain separated. Assessment thus follows the direction in which the model would be used and keeps future outcomes out of both model fitting and historical contexts [36, 37, 38, 54].

2.6. Reference baselines and current-input models

With ℓ_i the line of operation i , the line-specific SEC reference estimates one energy intensity per line from training totals:

$$\widehat{s}_\ell = \frac{\sum_{i \in \mathcal{T}: \ell_i = \ell} y_i}{\sum_{i \in \mathcal{T}: \ell_i = \ell} m_i}, \quad \widehat{y}_i^{\text{SEC}} = m_i \widehat{s}_{\ell_i}. \quad (3)$$

A line-specific median-energy reference and a recent-SEC reference (median of the last five completed same-line energy-to-mass ratios, updated as outcomes arrive during validation) complete the set of operational references. With z_i the transformed current-record inputs, Ridge regression estimates an unpenalized




Figure 2: Chronological allocation by pelleting line and partition. (a) Daily record counts for the three lines; background shading marks the development, calibration and frozen-test partitions, and dashed lines mark the development–calibration boundary on 22 February 2025 and the start of testing on 6 March 2025. Each record is counted by its earliest recorded stage start. (b) Dates, candidate records, exclusions and retained sample sizes of each partition.

intercept and regularized coefficients,

$$(\hat{b}, \hat{\beta}) = \arg \min_{b, \beta} \left\{ \sum_{i \in \mathcal{T}} (y_i - b - z_i^\top \beta)^2 + 10 \|\beta\|_2^2 \right\}, \quad \hat{y}_i = \max(0, \hat{b} + z_i^\top \hat{\beta}), \quad (4)$$

where the non-negative truncation prevents negative energy estimates without imposing a process-energy balance. ExtraTrees (250 trees, minimum leaf size five), histogram gradient boosting (200 iterations, seven leaves, learning rate 0.05, minimum leaf size 20, regularization 10) and a multilayer perceptron (hidden layers of 32 and 16 units, Adam, regularization 0.1, at most 800 iterations, seeds 11, 29 and 47 averaged) use the same inputs with configurations held constant across

validation blocks; the tree methods use seed 2026. These are the implementations evaluated here, not an exhaustive optimization of each family. When Ridge and gradient boosting gave practically the same development MAE, Ridge was chosen for parsimony, stability and record-level traceability. The selected model and its interval calibration were frozen before the test records were evaluated.

For interpretation, the 22 design columns are assigned to eight groups (dosed mass, batch count, hour, weekday, line, subfamily, presentation, bag/bulk), so that the untruncated estimate and the contribution of group g are

$$f_i^{\text{Ridge}} = \hat{b} + \sum_{g \in \mathcal{G}} C_{ig}, \quad C_{ig} = \sum_{k \in g} z_{ik} \hat{\beta}_k, \quad \hat{y}_i = \max(0, f_i^{\text{Ridge}}). \quad (5)$$

The grouped contributions preserve the arithmetic of the fitted model; they describe how the encoded variables build an estimate, not separate physical energy uses or causal effects [55, 56, 57].

2.7. Short-context temporal models on irregular operational events

The temporal comparison asks whether histories of completed same-line operations add information beyond the current-record covariates. Each eligible target receives the 16 most recent completed records from the same line and validation interval,

$$\mathcal{H}_i^{(16)} = \text{last}_{16, (a_j, j)} \{j : \ell_j = \ell_i, t_j \in \mathcal{P}, a_j < t_i, a_j < \sup \mathcal{P}\}, \quad (6)$$

where a_j is the recorded completion time and ties are broken by source row; an operation that started earlier but was not complete at the target origin cannot enter the context. The window retains 584 of the 800 development targets (149, 257 and 178 from Pellet 1, 2 and 3 over 21 origin dates). Contexts restart at the weekly validation boundaries; completed outcomes may inform later targets within the week but never update fitted parameters. Figure 3 shows the design and one observed context.

Two Ridge controls separate historical information from model capacity: one uses the current covariates only, the other appends the 16 historical covariate vectors and their energies, keeping the penalty of 10. TFT and N-HITS receive the same information through their NeuralForecast implementations [58, 59]; TFT uses hidden size 32, four heads and dropout 0.1, N-HITS three blocks of two 64-unit layers with pooling 2, 2, 1. Each uses seeds 11, 29 and 47 with an inner validation week before each outer week for epoch selection (Adam at 10^{-3} , weight decay 10^{-4} , batch 64, gradient clipping 1, at most 40 epochs, patience 8); each seed is refitted on

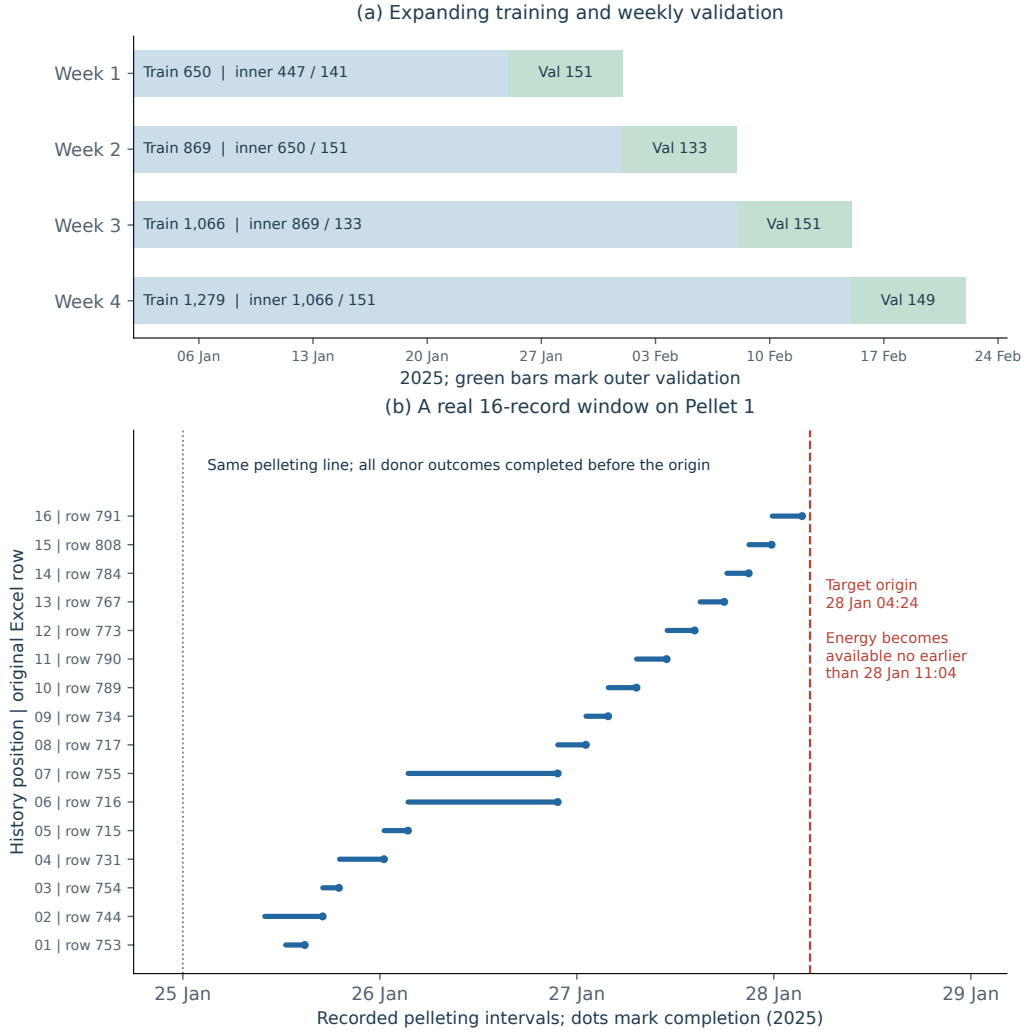


Figure 3: Short-context temporal evaluation. (a) Expanding training blocks and weekly validation blocks; green bars mark the outer validation weeks, and each training bar reports the earlier training / preceding-week validation counts used for epoch selection. (b) An observed context of 16 completed same-line operations for target row 823 on Pellet 1, ordered by completion; segments span recorded pelleting start and finish, dots mark completion, and every donor outcome precedes the target origin on 28 January at 04:24.

the full outer-training block for its selected epochs and the three clipped predictions are averaged. TimesFM 2.5 runs zero-shot with its frozen 200M checkpoint, alone and with its XReg adapter (ridge 10, one call per origin so that no query pools labels from another query’s future) [25, 60, 61]; Chronos-2 runs with energy-only and native covariate inputs, median forecast, cross-query learning disabled [62, 63]. All variants are evaluated on the same 584 targets, so the comparison concerns these short-context implementations on irregular pelleting events, not the model families in general [64, 65].

2.8. Post-hoc extensions

Three extensions turn patterns seen during development into hypotheses for later periods. H1 lets the mass association vary by line and H2 by product subfamily through the untruncated predictor

$$f_i = b + z_i^\top \beta + z_{m,i} \sum_{c \in C} \gamma_c \mathbf{1}(g_i = c), \quad (7)$$

with grouping variable g_i , the same penalty and training-only transformations. H3 corrects the frozen estimate with up to 20 recent same-line residuals $e_j = y_j - \hat{y}_j^{\text{Ridge}}$ of completed operations, ordered by completion and source row; with $n_i = |\mathcal{H}_i|$,

$$\delta_i = \begin{cases} \frac{n_i}{n_i + 10} \frac{1}{n_i} \sum_{j \in \mathcal{H}_i} e_j, & n_i \geq 5, \\ 0, & n_i < 5, \end{cases} \quad \hat{y}_i^{\text{H3}} = \max(0, \hat{y}_i^{\text{Ridge}} + \delta_i). \quad (8)$$

H1–H3 are exploratory development analyses; they do not replace the frozen model.

2.9. Residual intervals and performance metrics

Absolute errors on the 380 calibration records calibrate the intervals of the frozen model. For miscoverage level α ,

$$q_\alpha = r(\lceil (n_{\text{cal}} + 1)(1 - \alpha) \rceil), \quad r_j = |y_j - \hat{y}_j|, \quad I_i^{1-\alpha} = [\max(0, \hat{y}_i - q_\alpha), \hat{y}_i + q_\alpha], \quad (9)$$

giving half-widths of 93.01 and 132.96 kWh at the 90% and 95% levels. They are called residual intervals because the industrial sequence is not assumed exchangeable; they are judged empirically by coverage, mean width and interval score, globally and by line. A post-hoc rolling policy recomputes the quantile from the latest 200 available absolute errors, starting from the calibration residuals and admitting a test residual only once the completion of its operation precedes the

target origin ($a_j < t_i$); it is an exploratory comparator, not a validated improvement [66, 67, 68].

Point performance is measured by

$$\text{MAE} = N^{-1} \sum_i |\hat{y}_i - y_i|, \quad \text{RMSE} = \sqrt{N^{-1} \sum_i (\hat{y}_i - y_i)^2}, \quad (10)$$

$$\text{Bias} = N^{-1} \sum_i (\hat{y}_i - y_i), \quad \text{WAPE} = \frac{\sum_i |\hat{y}_i - y_i|}{\sum_i y_i}, \quad (11)$$

with MAE as the selection criterion [69]. Intervals are assessed by coverage, mean width and the interval score [51],

$$\text{IS}_\alpha(L, U; y) = (U - L) + \frac{2}{\alpha}(L - y)\mathbf{1}(y < L) + \frac{2}{\alpha}(y - U)\mathbf{1}(y > U), \quad (12)$$

and each completed operation yields a signed excess $r_i^{\text{excess}} = y_i - \hat{y}_i$ and an out-of-interval flag $B_i = \mathbf{1}\{y_i < L_i \text{ or } y_i > U_i\}$ that mark candidates for review without identifying the physical cause. Paired differences use a circular block bootstrap over observed origin dates (blocks of three dates, all lines of a date together; 3,000 replicates for the test comparison, 5,000 for the exploratory ones), with Bonferroni-adjusted 98.33% intervals for H1–H3 and 99.29% intervals for the seven temporal contrasts alongside the marginal 95% intervals.

2.10. Computational implementation

Ridge and the native JAX XReg solver ran on CPU; neural computations used float32 on an NVIDIA RTX PRO 5000 Blackwell GPU (48 GB) with TF32 disabled. TFT and N-HITS were trained locally, whereas TimesFM and Chronos weights remained frozen. CPU–GPU numerical comparisons are reported in Appendix [Appendix A](#).

3. Results

3.1. Production variability across lines and periods

Figure 4 summarizes production and electricity on the three lines: panels (a)–(c) contain the 2,742 records with positive pelleting energy, and panels (d)–(f) the mass, dosing-batch and product-family distributions of all 2,745 records. Electricity rises with dosed mass, yet records of similar mass remain widely dispersed, and the distributions of mass, batch count and product family differ among lines and

partitions. Both facts argue for a production-adjusted baseline rather than a single plant-wide intensity. Line-level differences in consumption must be read together with the corresponding production allocation; they are not controlled comparisons and say nothing about intrinsic mechanical efficiency.




Figure 4: Production and electricity across pelleting lines. (a–c) Dosed mass and pelleting electricity for the 2,742 positive-energy records, with line-specific medians annotated. (d–f) Mass distributions, dosing-batch frequencies and product-family mix of all 2,745 source records; values inside the product bars are record counts. Points are individual operations and are not uniformly spaced in time.

3.2. Baseline selection and frozen chronological test

On the 800 development targets (Table 4), Ridge and histogram gradient boosting were practically tied at 51.48 and 51.50 kWh MAE, with ExtraTrees

close behind at 52.31 kWh; the MLP and the three references had larger errors, and the recent-SEC estimator did worse than the static line-specific SEC despite using completed validation outcomes. Ridge was selected for comparable accuracy, lower structural complexity and an exact additive decomposition, not because it was shown to beat gradient boosting.

Table 4: Chronological development performance on the same 800 records. The recent-SEC estimator was updated as completed outcomes became available during validation; all other models kept fixed parameters within each validation block.

Method	MAE (kWh)	RMSE (kWh)	Bias (kWh)	WAPE (%)
Ridge	51.48	72.31	−3.14	19.88
Gradient boosting	51.50	73.46	−5.72	19.88
ExtraTrees	52.31	75.37	−4.89	20.20
MLP, three-seed mean	56.64	81.36	−1.47	21.87
Line-specific SEC	59.69	83.92	−6.59	23.05
Recent SEC, last five records	73.51	109.55	4.29	28.38
Line-specific median	82.61	112.78	−14.09	31.89

Fitted on the 1,542 eligible development records and frozen, Ridge reduced test MAE from 61.29 to 52.45 kWh (Table 5): a paired reduction of 8.84 kWh per record, or 14.42%, with a 95% block-bootstrap interval of [6.08, 11.50] kWh. RMSE fell from 92.05 to 80.42 kWh and global WAPE from 22.34% to 19.11%. These are reductions in prediction error, not in plant electricity use. Removing the five test records longer than 24 hours gave 52.32 kWh for Ridge and 61.15 kWh for SEC, so those records do not explain the advantage.

The improvement appeared on every line (Figure 5) but the remaining errors were uneven. Pellet 2 and Pellet 3 were the most consistent, with Ridge MAE of 43.36 and 53.36 kWh. Pellet 1 also benefited, its MAE falling from 72.31 to 66.37 kWh and its bias from −64.03 to −53.85 kWh, but the remaining underestimation means the frozen model was not operationally calibrated for that line during the test period. This localized shortfall does not invalidate the baseline; it does mean that a common alerting criterion should not be applied across lines without line-specific calibration.

Table 5: Frozen chronological test performance overall and by pelleting line. Global WAPE is 19.11% for Ridge and 22.34% for line-specific SEC.

Population	Method	N	MAE (kWh)	RMSE (kWh)	Bias (kWh)
All lines	Ridge	798	52.45	80.42	-21.89
All lines	Line-specific SEC	798	61.29	92.05	-24.57
Pellet 1	Ridge	200	66.37	105.14	-53.85
Pellet 1	Line-specific SEC	200	72.31	113.90	-64.03
Pellet 2	Ridge	333	43.36	61.26	-10.85
Pellet 2	Line-specific SEC	333	49.29	73.24	-8.66
Pellet 3	Ridge	265	53.36	80.10	-11.63
Pellet 3	Line-specific SEC	265	68.05	94.79	-14.79




Figure 5: Frozen Ridge performance in the chronological test. (a–c) Observed versus predicted electricity on identical axes, with the equality line dashed and MAE, RMSE and bias in kWh. (d–f) Signed errors in chronological order within each line, with zero dashed and the mean signed error of each line in red. The 200, 333 and 265 records of Pellet 1, 2 and 3 make up all 798 test observations.

3.3. Predictive value of completed operating histories

The short-context comparison uses the 584 development records for which 16 completed same-line operations were available within the weekly interval; the other 216 did not meet the requirement, so these results are not comparable with the 800-record ranking. Current-input Ridge reached 53.55 kWh MAE and TFT 53.65 kWh (Table 6); the paired difference (Ridge minus TFT) was -0.11 kWh with a 95% interval of $[-2.20, 2.32]$ kWh, which includes effects in both directions and gives no support for an improvement from TFT. Giving Ridge the full 16-record history raised its MAE to 63.95 kWh, and N-HiTS reached 65.28 kWh. Chronos-2 did better with covariates than with energy alone, but both versions stayed well above the Ridge control, as did the energy-only TimesFM variant (Figure 6).

Table 6: Short-context model performance on the 584 matched development records. TFT and N-HiTS are three-seed prediction means; foundation-model weights remain frozen.

Model group	Method	MAE	RMSE	Bias	WAPE (%)
Ridge controls	Current inputs	53.55	77.01	-4.72	20.37
	Full 16-record window	63.95	85.95	-1.42	24.32
Local temporal	TFT	53.65	79.11	-11.23	20.41
	N-HiTS	65.28	92.56	-3.05	24.83
Frozen pretrained	Chronos-2 + covariates	85.53	120.82	-8.11	32.53
	TimesFM 2.5, energy only	96.40	130.72	-9.22	36.67
	Chronos-2, energy only	100.02	136.47	-11.60	38.04
	TimesFM 2.5 + XReg	112.12	644.61	28.03	42.65

MAE, RMSE and bias are in kWh per record.

The TimesFM XReg variant produced three extreme Pellet 3 predictions of about 2,753, 9,239 and 12,203 kWh for observed values of 185.4, 142.4 and 70.9 kWh: in those contexts the dosed mass of the 16 donors was almost constant, so the standardized current mass lay roughly 1,100–1,467 historical standard deviations from the local mean and the adapter extrapolated accordingly. An independent float64 reconstruction reproduced the behaviour, and the three predictions are retained in every metric. The diagnosis concerns this short-context XReg pipeline, not TimesFM with longer contexts or other scaling policies. Under the tested representation, historical inputs did not improve prediction over current-record covariates; results by week, line and seed and the paired intervals of every contrast are given in Supplementary Tables S12–S16 and Figs. S3–S4.

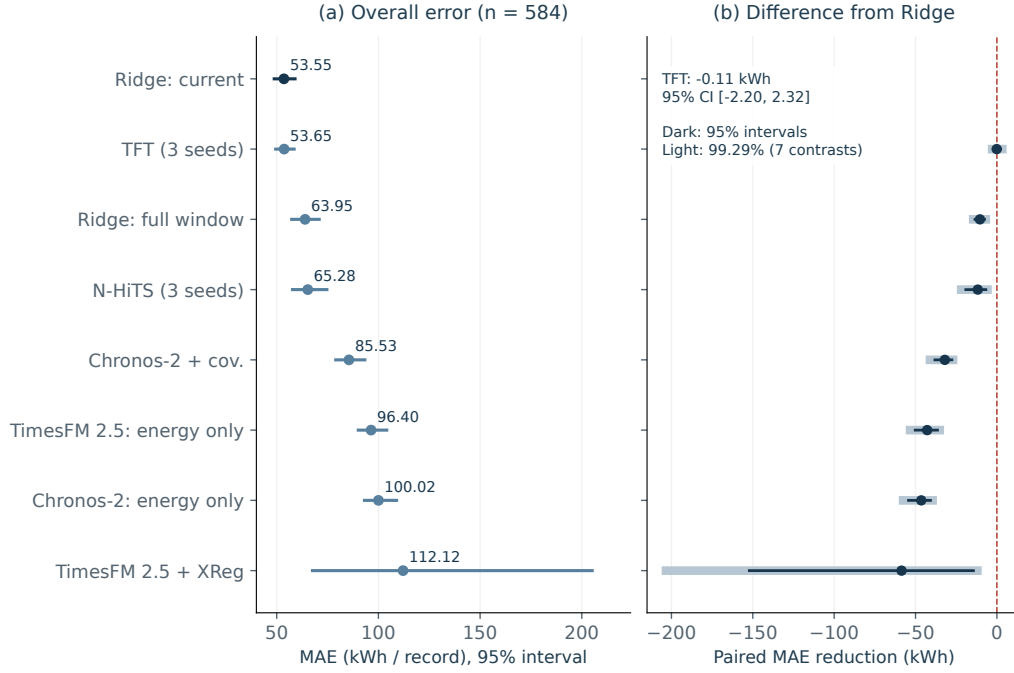



Figure 6: Short-context comparison on the 584 matched development records. (a) MAE with 95% block-bootstrap intervals. (b) Paired MAE reduction relative to current-input Ridge, so positive values favour the alternative; pale intervals are adjusted for seven comparisons and dark intervals are unadjusted 95% intervals. The wide TimesFM XReg interval is shown in full.

3.4. Additive decomposition of the fitted predictor

Figure 7 shows what the baseline returns for one operation: panel (a) the transformation of recorded production conditions into an expected electricity value, a residual interval and a signed deviation, and panel (b) the output for source row 2080, a Pellet 2 operation with 18.002 t in five dosing batches. The model estimated 255.90 kWh; the recorded value was 267.30 kWh, a residual of 11.40 kWh inside the fixed 95% interval of [122.94, 388.85] kWh. The record was chosen because its energy is close to the median test value, not because its error is typical.

Applying Equation 5 to every frozen estimate reproduced all 798 archived test values to within 1.2×10^{-13} kWh, and none required clipping. Figure 8 separates the fitted coefficients from record-specific contributions: panel (a) gives the numerical coefficients, (b) the categorical coefficients under the fitted coding, (c) the distributions of the grouped contributions and (d) the decomposition of source row 2080, whose intercept and eight grouped contributions sum exactly to

a From production conditions to an energy baseline



b Reading the output: an actual operation

Pellet 2 | 10 March 2025

18.002 t; 5 dosing batches
Swine-finishing product
Pelleted; bagged
Origin 08:31; completion 11:21

Source row 2080




Figure 7: Production-adjusted baseline output. (a) Recorded production conditions are transformed into an expected electricity value, residual intervals and a signed deviation; residuals from the later calibration period set the interval limits. (b) A real Pellet 2 operation (source row 2080): the light segment is the fixed 95% interval, the dark segment the 90% interval, the circle the estimate and the diamond the observed energy. An excess above the estimate is a model residual, not a measured saving or a confirmed fault.

255.90 kWh. After undoing the standardization, the mass coefficient is 13.30 kWh per additional tonne and the batch coefficient 9.52 kWh per additional batch, conditional on the other encoded predictors (all 22 coefficients and conversion factors are in Supplementary Tables S6–S7).

These values are coefficients of a fitted linear predictor, not separately identified physical requirements: mass and batch count are correlated, and line and product allocation were not randomized. The intercept of 288.09 kWh is the algebraic reference of the coding, not a standby consumption; categorical coefficients depend




Figure 8: Additive decomposition of the frozen Ridge predictor. (a) Numerical coefficients in original input units. (b) All 16 categorical coefficients under the fitted coding. (c) Grouped contributions for every test record, with medians in red and interquartile ranges in orange. (d) Arithmetic decomposition of source row 2080. The reference combines standardized numerical means with all-zero categorical indicators and need not describe a feasible operating condition; contributions are conditional model terms, not causal effects.

on the one-hot representation; and small line coefficients do not demonstrate equal intrinsic efficiency. What the exact equality provides is record-level traceability.

3.5. *Post-hoc extensions*

On the 800 development records (Table 7), letting the mass association vary by product subfamily reduced MAE by 1.44 kWh (2.80%) and the recent-residual correction by 1.71 kWh (3.33%), both with multiplicity-adjusted intervals above zero; the line-dependent mass interaction reduced MAE by only 0.22 kWh, with an interval crossing zero. The gains varied across lines and weeks (Supplementary Fig. S1 and Tables S8–S9). Because H1–H3 were motivated by development patterns, they are hypothesis-generating and need a new chronological period; they do not replace the frozen specification.

Table 7: Post-hoc exploratory point-model extensions on the 800 development records. Positive reductions favour the extension; the adjusted intervals cover these three comparisons only.

Variant	MAE	MAE reduction	Adjusted 98.33% interval
H1: mass by pelleting line	51.26	0.22	[−0.08, 0.58]
H2: mass by subfamily	50.04	1.44	[0.55, 2.48]
H3: recent residuals	49.77	1.71	[0.24, 3.16]

All quantities are in kWh per record.

3.6. *Residual intervals and records flagged for review*

The fixed 90% and 95% intervals covered 85.96% and 93.23% of the 798 test records (Figure 9) with mean widths of 183.63 and 258.85 kWh; their block-bootstrap 95% coverage intervals, [81.86, 90.24]% and [90.53, 95.87]%, make global coverage broadly compatible with the nominal levels. Line-specific coverage differed: at the 95% level it was 90.00% on Pellet 1 (20 records outside), 96.40% on Pellet 2 (12) and 91.70% on Pellet 3 (22). In total 54 records fell outside the 95% intervals, 45 above and nine below: about 6.77 flagged records per 100 operations. They are candidates for technical review because the calibrated baseline did not represent their energy well; without plant evidence they should not be called anomalies, inefficient operations or avoidable consumption.

The post-hoc rolling policy, which updates the quantiles from the most recent 200 completed residuals, raised coverage to 89.47% and 94.49% but widened the intervals by 19.66% and 18.85% while improving the interval score by only 1.12% and 0.23% (Table 8); most of the additional coverage came from width, and the

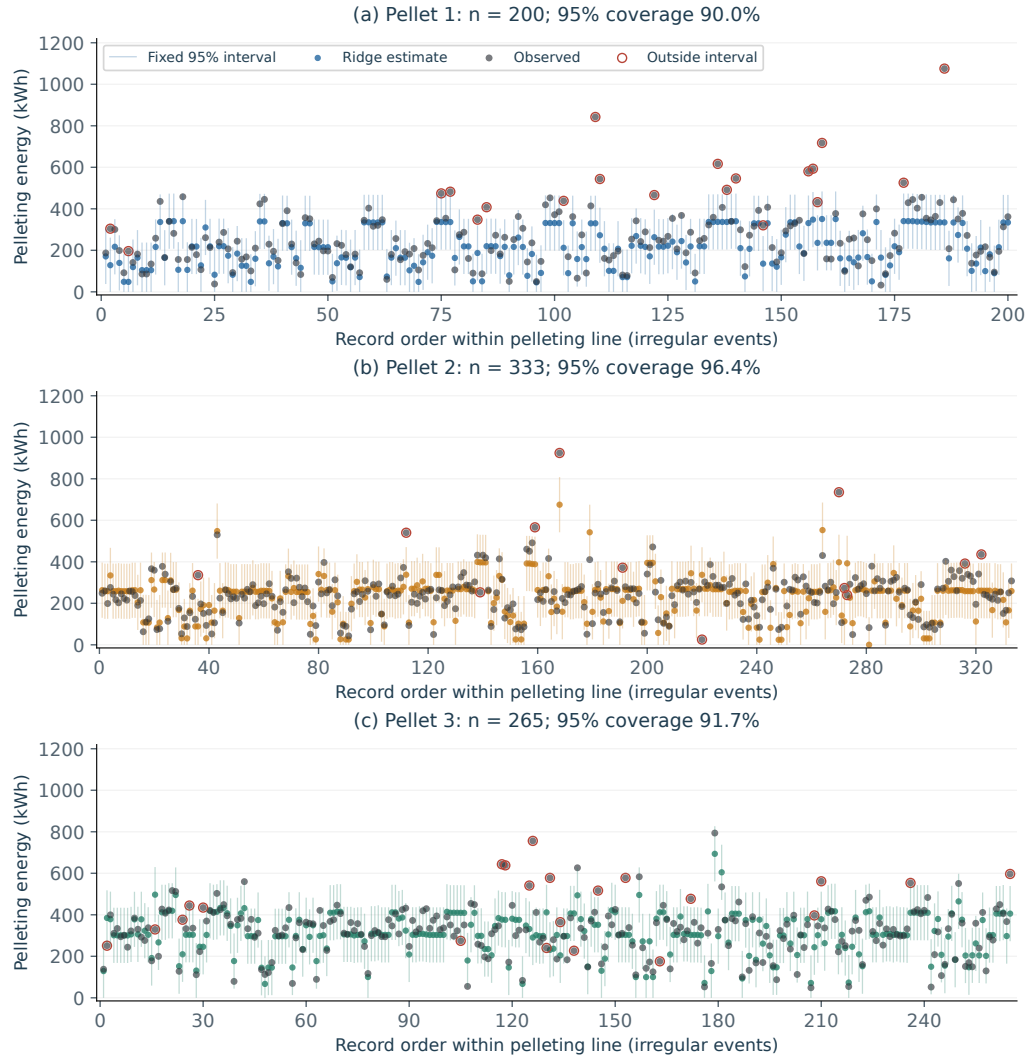


Figure 9: Frozen Ridge estimates and fixed 95% residual intervals for every test record, in chronological order within each pelleting line. Black points are observed electricity, coloured points the frozen estimate and vertical lines the intervals; red circles mark the 54 outcomes outside their intervals. Horizontal position is record order, so equal spacing does not imply equal elapsed time.

number of records outside the 95% interval fell from 54 to 44. Because the policy was defined after the test results had been inspected, it is a calibration diagnostic, not a validated replacement for the fixed intervals; coverage, width and interval score by line and week are given in Supplementary Tables S10–S11 and Fig. S2.

Table 8: Fixed and post-hoc rolling residual intervals on the same 798 frozen point predictions. The rolling policy has not been validated on an independent period.

Policy	Nominal level	Coverage (%)	Mean width	Interval score
Fixed	90%	85.96	183.63	381.79
Rolling, 200 residuals	90%	89.47	219.74	377.51
Fixed	95%	93.23	258.85	497.64
Rolling, 200 residuals	95%	94.49	307.64	496.48

Width and interval score are in kWh per record.

4. Discussion

4.1. Production adjustment against intensity-based references

During development, Ridge produced lower prediction errors than the line-specific median, recent SEC, MLP, ExtraTrees and line-specific SEC references and performed like histogram gradient boosting; in the frozen chronological test its MAE was 14.42% below line-specific SEC. This agrees with energy-baseline research that recommends adjusting consumption for the relevant operating conditions rather than relying on aggregate or intensity-based indicators [12, 13, 14, 15], and it extends the pelleting studies that link energy demand to production volume, formulation, conditioning, equipment and operating conditions [10, 7, 8, 9, 11]. The contribution is a chronologically evaluated baseline that combines production mass, product characteristics and line identity for heterogeneous completed operations, which is the kind of normalization that ISO 50006 and the IPMVP expect of an energy baseline [2, 3]. The improvement is in prediction accuracy: it is neither a reduction in actual energy use nor an estimate of achievable savings, and Ridge was chosen for its balance of accuracy, parsimony and traceability rather than for universal superiority over nonlinear models.

4.2. Line-specific behaviour

Production adjustment helped on all three lines. Pellet 2 and Pellet 3 showed the most consistent agreement between predictions and observations and are the

strongest basis for prospective validation. Pellet 1 also improved on line-specific SEC, but its remaining negative bias shows that the frozen model was not operationally calibrated for that line. This localized requirement does not undermine the common modelling framework; it shows that aggregate performance can hide line-specific deficiencies. Differences among lines may reflect equipment configuration, production allocation, maintenance, metering or unrecorded operating conditions and should not be read as intrinsic differences in energy efficiency.

4.3. Operating histories, temporal models and foundation models

Under the tested 16-record, same-line, within-week representation, completed-operation histories did not improve on current-record covariates. Dosed mass, product characteristics and line identity may already capture much of the repeatable structure, and irregular intervals weaken the meaning of sequential proximity. The conclusion is confined to the evaluated short-context representation and does not contradict the value of temporal models where sequences carry informative dependence [23, 24]; likewise, the TimesFM and Chronos results characterize the tested interfaces, not the general capabilities of foundation models [25, 26, 62]. The XReg extrapolation and the device-dependent normalization sensitivity of Chronos-2 (Appendix [Appendix A](#)) show that covariate adaptation and normalization shape the final industrial prediction as much as the pretrained weights do, and that a model’s ability to accept the available covariates is no evidence that its pipeline is stable for this kind of record. Longer contexts, explicit elapsed time, operating-state information, maintenance events and regularly sampled measurements would allow a fairer assessment of temporal information in industrial pelleting [64, 65].

4.4. Reading the additive decomposition

Ridge gives an exact additive decomposition of its fitted linear predictor and therefore a direct, reproducible audit trail. Its coefficients and record-specific contributions nevertheless depend on scaling, categorical coding, reference categories, feature construction and correlations among predictors [55, 56, 57]: they describe the fitted prediction function, do not identify causal effects and would change under alternative coding choices. Line and product terms may also absorb unrecorded differences in formulation, equipment condition, production allocation or operating practice. The small exploratory gains from product-dependent mass effects and residual correction indicate where the next refinement may come from, but they are hypotheses for a new period, not reasons to replace the frozen model on the strength of an already examined test.

4.5. Uncertainty and the meaning of a flagged record

The residual intervals showed that satisfactory marginal coverage does not guarantee comparable line-specific calibration [49, 47, 48]. Pellet 2 had the most stable coverage, whereas Pellet 1 and, to a lesser extent, Pellet 3 need further assessment. Rolling calibration raised coverage mainly by widening the intervals, which confirms that coverage must be read together with interval width, bias and alert frequency [50, 51]. An interval exceedance is therefore a candidate for technical review, not a confirmed anomaly, inefficiency, equipment failure or energy-saving opportunity.

4.6. Scope of use and digital-twin integration

The evidence supports retrospective assessment after each recorded operation closes: an expected electricity value, an empirically calibrated residual interval and a traceable deviation that identifies candidates for review. Prospective supervisory alerts require validation over a new consecutive production period with prespecified line-level criteria and, for Pellet 1, additional recalibration or predictors. Automatic control and modification of process set-points remain outside the evidence, because the observational model does not estimate causal responses. A digital-twin application should therefore begin with validated monitoring and decision support [52], storing the input and model version of each issued estimate and attaching the measured outcome when it becomes available, so that an earlier forecast and a later production-adjusted assessment remain distinguishable.

4.7. Strengths and limitations

The study draws its strength from real industrial records, chronological separation, operational comparators, line-specific diagnostics, residual intervals and a reproducible decomposition. Its transferability is limited by the absence of a verified production-order identifier, retained temporal overlaps, unverified meter boundaries and undocumented predictor-availability times. The data cover three lines of one plant over about three months, use product categories rather than complete formulations, and lack detailed maintenance and operating-state variables. The intervals mainly assess marginal coverage, several extensions were post hoc, and the operational records cannot be released publicly, although they are available from the corresponding author on reasonable request. The contribution therefore lies in the evaluation framework—production adjustment, chronological testing, line-specific assessment, uncertainty quantification and algebraic traceability—rather than in a universally deployable fitted model.

5. Conclusions

We set out to establish whether an explainable, production-adjusted baseline predicts the electricity of a completed feed-pelleting operation better than the line-specific SEC reference used in practice, and whether operating histories, temporal models or pretrained foundation models add anything to it. On 798 real records evaluated after the model had been frozen, the production-adjusted Ridge baseline reduced MAE by 14.42% relative to line-specific SEC on all three lines, while every estimate remained exactly decomposable into its inputs. The improvement concerns the estimation of expected electricity under the recorded production conditions; it is not a reduction in plant energy use or evidence of achievable savings.

Under the tested 16-record, same-line, within-week representation, completed operating histories did not improve on current-record covariates, and neither TFT, N-HiTS nor the tested TimesFM 2.5 and Chronos-2 configurations beat the current-input control on identical records. Production adjustment helped most consistently on Pellet 2 and Pellet 3; Pellet 1 improved on SEC but its systematic underestimation calls for line-specific recalibration or additional predictors before prospective use.

The evidence supports retrospective assessment at the closure of recorded operations, where the baseline provides an expected value, an empirically calibrated residual interval and a traceable deviation for identifying candidates for technical review. Prospective supervisory use requires validation on a new consecutive production period with prespecified line-level criteria and verified record identity, measurement boundaries, predictor availability and calibration rules. Until then, automatic control, modification of process set-points and operational digital-twin integration remain outside the demonstrated scope.

CRediT authorship contribution statement

Vladimir Sousa Santos: Conceptualization, Methodology, Investigation, Writing – original draft, Project administration, Funding acquisition. **Juan José Cabello Eras:** Conceptualization, Supervision, Resources, Writing – review & editing. **Eder Elías Herrera Rodiño:** Investigation, Data curation, Resources, Writing – review & editing. **Francisco Manuel Arrabal-Campos:** Methodology, Software, Formal analysis, Data curation, Validation, Visualization, Writing – original draft, Writing – review & editing.

Funding

The authors gratefully acknowledge the Asociación Universitaria Iberoamericana de Postgrado (AUIP) and the Consejería de Universidad, Investigación e Innovación of the Junta de Andalucía for the support received through the Mobility Grants programme for Postdoctoral Stays at Andalusian Universities (2026 call, grant PMP-2026-06-045), which made possible the research stay at the University of Almería.

Declaration of competing interest

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Acknowledgements

The authors thank the Corporación Universidad de la Costa and the University of Almería for their institutional support, and the Master’s Programme in Mechanical Engineering of the Universidad de Córdoba for its participation in the study and the management of the data used.

Supplementary material

Supplementary material associated with this article (the full source-variable dictionary, descriptive statistics, the complete coefficient table of the frozen Ridge model, bootstrap and calibration details, residual-interval diagnostics by line and week, stratified and per-seed results of the temporal benchmark, CPU–GPU computational diagnostics and four additional figures) is provided as a separate document.

Data availability

The code, frozen model configurations, archived aggregate results, figure scripts and LaTeX sources are openly available in the GitHub repository <https://github.com/fmarrabal/pelleting-energy-baseline-real-data> (release v4-submission). The industrial production and electricity records, and the record-level tables derived from them, belong to the plant and are not publicly available; they can be obtained from the corresponding author (fmarrabal@ual.es) on reasonable request, subject

to the data owner’s agreement. Pretrained checkpoints are identified by repository and revision and are not redistributed.

Appendix A. Computational and numerical reproducibility

The GPU experiment used Python 3.13.13, PyTorch 2.14.0 with CUDA 13.0, NeuralForecast 3.2.2, the TimesFM package 3.0.2, Chronos-forecasting 2.3.2 and scikit-learn 1.9.0. The pretrained checkpoint revisions were 1d952420fba87f3c6dee4f240de0f1a0fbc790e3 for TimesFM 2.5 and 29ec3766d36d6f73f0696f85560a422f50e8498c for Chronos-2. TimesFM received 16 observed context records, padded to a maximum context of 32, with a configured maximum horizon of 32; only its first forecast position was evaluated. Chronos used float32, disabled cross-query learning and returned the median forecast. Model configurations, checkpoints and record-level outputs are retained in the repository; pretrained weights are identified by revision rather than redistributed.

All 128,896 logged donor relationships passed the checks for same-pelleting-line membership, partition membership and completion before the target origin; the total includes repeated relationships across nested and expanding blocks. Source-data, frozen-model and final-prediction fingerprints remained unchanged, and the shared CPU and GPU data, window and encoding functions were checked for consistency. These checks concern the recorded completion-time proxy, since actual publication delays are not available in the source.

Within the GPU run, the 24 saved neural checkpoints were reloaded and compared with their stored inference, together with four Chronos covariate batch-size checks. All 28 checks passed the recorded tolerances with a maximum difference of 7.63×10^{-5} kWh. Fresh CPU and GPU neural fits were less closely aligned, as optimization can follow different trajectories even with the same integer seed: TFT’s CPU and GPU MAEs were 53.25 and 53.65 kWh, and neither established an advantage over the matched Ridge control under the paired intervals. Figure A.10 shows all three GPU seeds and their ensemble.

For frozen Chronos-2 with covariates, the median absolute CPU–GPU prediction difference was 0.86 kWh and the maximum 57.04 kWh, and its overall MAE moved from 85.65 to 85.53 kWh. Maximum prediction differences for both frozen TimesFM variants and for univariate Chronos stayed below 1.9×10^{-4} kWh. To trace the covariate sensitivity, the installed Chronos preprocessing and instance normalization were applied to the same 584 contexts, giving 6,424 target/covariate channels. Among the 1,097 constant historical channels, 694 showed a CPU–GPU

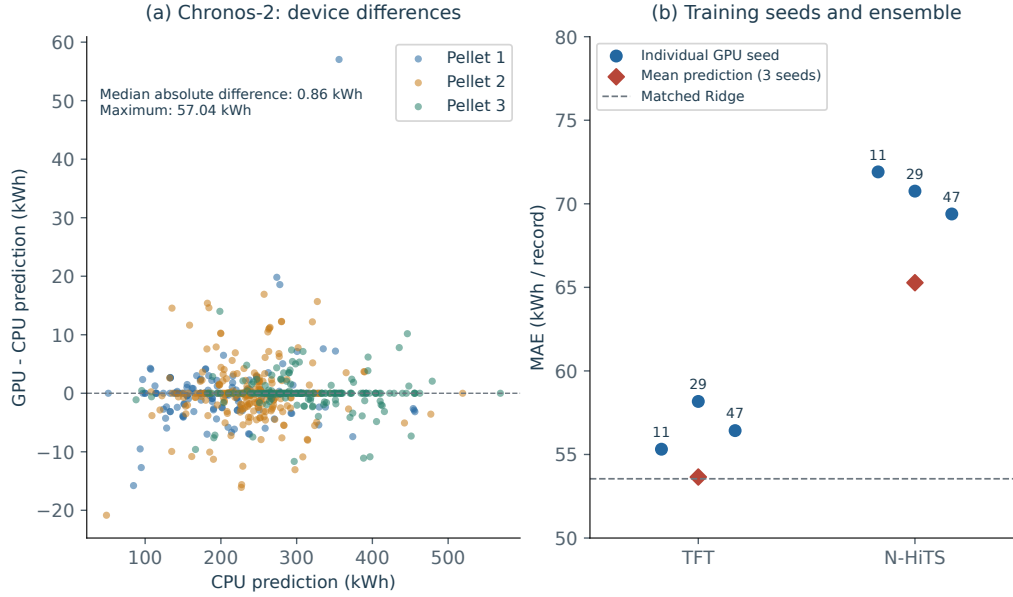


Figure A.10: Numerical and training variability. (a) Every paired CPU/GPU Chronos-2 covariate output as a signed difference against the CPU prediction. (b) All three GPU seeds (11, 29 and 47) and their prediction-mean ensemble for TFT and N-HITS, with the matched Ridge reference; averaging predictions is not the same as averaging the three MAEs. These diagnostics accompany the primary comparison without selecting a preferred seed or device.

difference above 0.1 in normalized context values, whereas no non-constant channel exceeded that threshold; the largest difference was 0.88137 in the transformed coordinates.

The implementation substitutes a small scale only when the calculated scale is exactly zero, so tiny round-off variances in constant channels can be amplified before the inverse-hyperbolic-sine transform. This mechanism is consistent with the observed sensitivity, although it is not an end-to-end causal demonstration for each prediction. The reported benchmark keeps the native preprocessing and original outputs; a corrective normalization policy would require a separately selected and evaluated experiment.

Declaration of generative AI and AI-assisted technologies in the writing process

During the preparation of this work the authors used Trinka as a language-editing assistant to improve the readability and English of the manuscript, and Claude Code as a coding assistant for writing and documenting the analysis scripts.

After using these tools, the authors reviewed and edited the content as needed and take full responsibility for the content of the published article.

References

- [1] International Organization for Standardization, ISO 50001:2018 Energy management systems — Requirements with guidance for use (2018).
- [2] International Organization for Standardization, ISO 50006:2023 Energy management systems — Evaluating energy performance using energy performance indicators and energy baselines (2023).
- [3] Efficiency Valuation Organization, International Performance Measurement and Verification Protocol: Core Concepts, EVO 10000-1:2022 (2022).
- [4] M. Compton, S. Willis, B. Rezaie, K. Humes, Food processing industry energy and water consumption in the pacific northwest, *Innovative Food Science & Emerging Technologies* 47 (2018) 371–383. [doi:10.1016/j.ifset.2018.04.001](https://doi.org/10.1016/j.ifset.2018.04.001).
- [5] A. Foster, T. Brown, J. Evans, Carbon emissions from refrigeration used in the uk food industry, *International Journal of Refrigeration* 150 (2023) 297–303. [doi:10.1016/j.ijrefrig.2023.01.022](https://doi.org/10.1016/j.ijrefrig.2023.01.022).
- [6] J. Park, T. Jin, S. Lee, J. Woo, Industrial electrification and efficiency: Decomposition evidence from the korean industrial sector, *Energies* 14 (16) (2021) 5120. [doi:10.3390/en14165120](https://doi.org/10.3390/en14165120).
- [7] M. Thomas, A. van der Poel, Physical quality of pelleted animal feed 1. Criteria for pellet quality, *Animal Feed Science and Technology* 61 (1-4) (1996) 89–112. [doi:10.1016/0377-8401\(96\)00949-2](https://doi.org/10.1016/0377-8401(96)00949-2).
- [8] M. Thomas, D. van Zuilichem, A. van der Poel, Physical quality of pelleted animal feed. 2. contribution of processes and its conditions, *Animal Feed Science and Technology* 64 (2-4) (1997) 173–192. [doi:10.1016/s0377-8401\(96\)01058-9](https://doi.org/10.1016/s0377-8401(96)01058-9).
- [9] M. Thomas, A. F. B. van der Poel, Fundamental factors in feed manufacturing: Towards a unifying conditioning/pelleting framework, *Animal Feed Science and Technology* 268 (2020) 114612. [doi:10.1016/j.anifeedsci.2020.114612](https://doi.org/10.1016/j.anifeedsci.2020.114612).

- [10] E. R. Skoch, K. C. Behnke, C. W. Deyoe, S. F. Binder, The effect of steam-conditioning rate on the pelleting process, *Animal Feed Science and Technology* 6 (1) (1981) 83–90. doi:[10.1016/0377-8401\(81\)90033-X](https://doi.org/10.1016/0377-8401(81)90033-X).
- [11] T. Bastiaansen, R. Benders, J. Dijksman, M. Thomas, W. Hendriks, S. de Vries, G. Bosch, Changes in thermomechanical properties of feed in relation to composition and their effect on pellet manufacturing, *Animal Feed Science and Technology* 295 (2023) 115540. doi:[10.1016/j.anifeedsci.2022.115540](https://doi.org/10.1016/j.anifeedsci.2022.115540).
- [12] J. Kelly Kissock, C. Eger, Measuring industrial energy savings, *Applied Energy* 85 (5) (2008) 347–361. doi:[10.1016/j.apenergy.2007.06.020](https://doi.org/10.1016/j.apenergy.2007.06.020).
- [13] J. Granderson, S. Touzani, C. Custodio, M. D. Sohn, D. Jump, S. Fernandes, Accuracy of automated measurement and verification (M&V) techniques for energy savings in commercial buildings, *Applied Energy* 173 (2016) 296–308. doi:[10.1016/j.apenergy.2016.04.049](https://doi.org/10.1016/j.apenergy.2016.04.049).
- [14] C. V. Gallagher, K. Bruton, K. Leahy, D. T. O’Sullivan, The suitability of machine learning to minimise uncertainty in the measurement and verification of energy savings, *Energy and Buildings* 158 (2018) 647–655. doi:[10.1016/j.enbuild.2017.10.041](https://doi.org/10.1016/j.enbuild.2017.10.041).
- [15] J. Cosgrove, J. Littlewood, P. Wilgeroth, Development of a framework of key performance indicators to identify reductions in energy consumption in a medical devices production facility, *International Journal of Ambient Energy* 39 (2) (2018) 202–210. doi:[10.1080/01430750.2017.1278718](https://doi.org/10.1080/01430750.2017.1278718).
- [16] G. C. Sánchez, J. P. Monteagudo Yanes, M. Montesino Pérez, J. L. Cabrera Sánchez, A. P. Padrón, D. Haeseldonckx, Efficiency in electromechanical drive motors and energy performance indicators for implementing a management system in balanced animal feed manufacturing, *Energy* 194 (2020) 116818. doi:[10.1016/j.energy.2019.116818](https://doi.org/10.1016/j.energy.2019.116818).
- [17] M. Wigger, P. Burggräf, F. Steinberg, A. Becher, B. Heinbach, Integrating artificial intelligence into energy management: A case study on energy consumption data analysis and forecasting in a German manufacturing company, *Energy and AI* 21 (2025) 100576. doi:[10.1016/j.egyai.2025.100576](https://doi.org/10.1016/j.egyai.2025.100576).

- [18] A. E. Hoerl, R. W. Kennard, Ridge Regression: Biased Estimation for Nonorthogonal Problems, *Technometrics* 12 (1) (1970) 55–67. doi:[10.1080/00401706.1970.10488634](https://doi.org/10.1080/00401706.1970.10488634).
- [19] J. H. Friedman, Greedy function approximation: A gradient boosting machine., *The Annals of Statistics* 29 (5) (2001) 1189–1232. doi:[10.1214/aos/1013203451](https://doi.org/10.1214/aos/1013203451).
- [20] H. Ekwaro-Osire, D. Bode, J.-H. Ohlendorf, K.-D. Thoben, Manufacturing process energy consumption modeling: A methodology to identify the most appropriate model, *Journal of Intelligent Manufacturing* 36 (2025) 5673–5693. doi:[10.1007/s10845-024-02514-z](https://doi.org/10.1007/s10845-024-02514-z).
- [21] C. Panjapornpon, S. Bardeeniz, M. A. Hussain, P. Chomchai, Explainable deep transfer learning for energy efficiency prediction based on uncertainty detection and identification, *Energy and AI* 12 (2023) 100224. doi:[10.1016/j.egyai.2022.100224](https://doi.org/10.1016/j.egyai.2022.100224).
- [22] F. Mohammed, A. Boumaiza, A. Sanfilippo, D. Perez-Astudillo, D. Bachour, A robust hybrid machine learning framework for short-term load forecasting: integrating multi-linear regression, long short-term memory, and feed-forward neural networks for enhanced accuracy and efficiency, *Energy and AI* 22 (2025) 100625. doi:[10.1016/j.egyai.2025.100625](https://doi.org/10.1016/j.egyai.2025.100625).
- [23] B. Lim, S. Ö. Arık, N. Loeff, T. Pfister, Temporal Fusion Transformers for interpretable multi-horizon time series forecasting, *International Journal of Forecasting* 37 (4) (2021) 1748–1764. doi:[10.1016/j.ijforecast.2021.03.012](https://doi.org/10.1016/j.ijforecast.2021.03.012).
- [24] C. Challu, K. G. Olivares, B. N. Oreshkin, F. Garza Ramirez, M. Mergenthaler Canseco, A. Dubrawski, NHITS: Neural Hierarchical Interpolation for Time Series Forecasting, *Proceedings of the AAAI Conference on Artificial Intelligence* 37 (6) (2023) 6989–6997. doi:[10.1609/aaai.v37i6.25854](https://doi.org/10.1609/aaai.v37i6.25854).
- [25] A. Das, W. Kong, R. Sen, Y. Zhou, [A decoder-only foundation model for time-series forecasting](#), in: *Proceedings of the 41st International Conference on Machine Learning*, Vol. 235 of *Proceedings of Machine Learning Research*, 2024, pp. 10148–10167.
URL <https://proceedings.mlr.press/v235/das24c.html>

- [26] A. F. Ansari, L. Stella, C. Turkmen, et al., [Chronos: Learning the Language of Time Series](#) (2024). [arXiv:2403.07815](#).
URL <https://arxiv.org/abs/2403.07815>
- [27] M. Obermeier, M. Pruckner, F. Haselbeck, A. Zeiselmaier, FETS Benchmark: Foundation Models Enable Scalable and Generalizable Energy Time Series Forecasting, *Energy and AI* 26 (2026) 100867. [doi:10.1016/j.egyai.2026.100867](#).
- [28] T. Miseta, A. Fodor, Á. Vathy-Fogarassy, Electricity price prediction with univariate Long Short-Term Memory and zero-shot foundation transformer models, *Energy and AI* 26 (2026) 100886. [doi:10.1016/j.egyai.2026.100886](#).
- [29] X. Lin, J. Ti, Benchmarking Chronos-2 for European wholesale electricity price forecasting: Evidence from foundation and conventional forecasting models, *Energy and AI* 26 (2026) 100914. [doi:10.1016/j.egyai.2026.100914](#).
- [30] L. Zhang, D. Diallo, C. Delpha, M. Benbouzid, Cross-dataset battery life forecasting with time-series foundation models: From zero-shot to prefix-adapted, *Energy and AI* 25 (2026) 100771. [doi:10.1016/j.egyai.2026.100771](#).
- [31] M. Beichter, N. Friederich, J. Pinter, D. Werling, K. Phipps, S. Beichter, O. Neumann, R. Mikut, V. Hagenmeyer, B. Heidrich, Decision-focused fine-tuning of time series foundation models for dispatchable feeder optimization, *Energy and AI* 21 (2025) 100533. [doi:10.1016/j.egyai.2025.100533](#).
- [32] A. Micallef, M. Apap, J. Licari, C. Spiteri Staines, Z. Xiao, A comparative framework for evaluating machine learning models in forecasting electricity demand for port microgrids, *Energy and AI* 20 (2025) 100494. [doi:10.1016/j.egyai.2025.100494](#).
- [33] B. Kavoosighafi, M. Eidenskog, W. Glad, K. Vrotsou, An empirical benchmark of deep time-series models for smart meter energy forecasting, *Energy and AI* 25 (2026) 100870. [doi:10.1016/j.egyai.2026.100870](#).
- [34] A. Neubauer, M. Yu, P. Babakhani, S. Brandt, M. Kriegel, Transfer learning and explainable AI for heating load forecasting: A large-scale benchmark

- with SHAP-based static features, *Energy and AI* 24 (2026) 100722. doi:[10.1016/j.egyai.2026.100722](https://doi.org/10.1016/j.egyai.2026.100722).
- [35] H. Hewamalage, K. Ackermann, C. Bergmeir, Forecast evaluation for data scientists: common pitfalls and best practices, *Data Mining and Knowledge Discovery* 37 (2) (2023) 788–832. doi:[10.1007/s10618-022-00894-5](https://doi.org/10.1007/s10618-022-00894-5).
 - [36] L. J. Tashman, Out-of-sample tests of forecasting accuracy: an analysis and review, *International Journal of Forecasting* 16 (4) (2000) 437–450. doi:[10.1016/s0169-2070\(00\)00065-0](https://doi.org/10.1016/s0169-2070(00)00065-0).
 - [37] V. Cerqueira, L. Torgo, I. Mozetič, Evaluating time series forecasting models: an empirical study on performance estimation methods, *Machine Learning* 109 (11) (2020) 1997–2028. doi:[10.1007/s10994-020-05910-7](https://doi.org/10.1007/s10994-020-05910-7).
 - [38] C. Bergmeir, R. J. Hyndman, B. Koo, A note on the validity of cross-validation for evaluating autoregressive time series prediction, *Computational Statistics & Data Analysis* 120 (2018) 70–83. doi:[10.1016/j.csda.2017.11.003](https://doi.org/10.1016/j.csda.2017.11.003).
 - [39] R. Machlev, L. Heistrene, M. Perl, K. Y. Levy, J. Belikov, S. Mannor, Y. Levron, Explainable Artificial Intelligence (XAI) techniques for energy and power systems: Review, challenges and opportunities, *Energy and AI* 9 (2022) 100169. doi:[10.1016/j.egyai.2022.100169](https://doi.org/10.1016/j.egyai.2022.100169).
 - [40] M. Saarela, V. Podgorelec, Recent applications of explainable ai (XAI): A systematic literature review, *Applied Sciences* 14 (19) (2024) 8884. doi:[10.3390/app14198884](https://doi.org/10.3390/app14198884).
 - [41] J. Cação, J. Santos, M. Antunes, Explainable ai for industrial fault diagnosis: A systematic review, *Journal of Industrial Information Integration* 47 (2025) 100905. doi:[10.1016/j.jii.2025.100905](https://doi.org/10.1016/j.jii.2025.100905).
 - [42] S. Letzgus, K.-R. Müller, An explainable AI framework for robust and transparent data-driven wind turbine power curve models, *Energy and AI* 15 (2024) 100328. doi:[10.1016/j.egyai.2023.100328](https://doi.org/10.1016/j.egyai.2023.100328).
 - [43] Y. Zhou, I. Aslanidou, M. Karlsson, K. Kyprianidis, An explainable AI model for power plant NOx emission control, *Energy and AI* 15 (2024) 100326. doi:[10.1016/j.egyai.2023.100326](https://doi.org/10.1016/j.egyai.2023.100326).

- [44] C. Samarajeewa, D. De Silva, M. Manic, N. Mills, H. Moraliyage, D. Alahakoon, An artificial intelligence framework for explainable drift detection in energy forecasting, *Energy and AI* 17 (2024) 100403. doi:[10.1016/j.egyai.2024.100403](https://doi.org/10.1016/j.egyai.2024.100403).
- [45] C. O'Connor, M. Bahloul, R. Rossi, S. Prestwich, A. Visentin, Conformal Prediction for electricity price forecasting in the day-ahead and real-time balancing market, *Energy and AI* 21 (2025) 100571. doi:[10.1016/j.egyai.2025.100571](https://doi.org/10.1016/j.egyai.2025.100571).
- [46] N. B. Nam, E. Ogliari, S. Leva, E. Pafumi, D. Alberti, M. Q. Duong, Comparative analysis of conformal prediction techniques and machine learning models for very short-term solar power forecasting, *Energy and AI* 21 (2025) 100573. doi:[10.1016/j.egyai.2025.100573](https://doi.org/10.1016/j.egyai.2025.100573).
- [47] A. N. Angelopoulos, S. Bates, [A Gentle Introduction to Conformal Prediction and Distribution-Free Uncertainty Quantification](https://arxiv.org/abs/2107.07511) (2021). arXiv: [2107.07511](https://arxiv.org/abs/2107.07511).
URL <https://arxiv.org/abs/2107.07511>
- [48] R. F. Barber, E. J. Candès, A. Ramdas, R. J. Tibshirani, Conformal prediction beyond exchangeability, *The Annals of Statistics* 51 (2) (2023) 816–845. doi:[10.1214/23-aos2276](https://doi.org/10.1214/23-aos2276).
- [49] R. Foygel Barber, E. J. Candès, A. Ramdas, R. J. Tibshirani, The limits of distribution-free conditional predictive inference, *Information and Inference: A Journal of the IMA* 10 (2) (2021) 455–482. doi:[10.1093/imaiai/iaaa017](https://doi.org/10.1093/imaiai/iaaa017).
- [50] T. Gneiting, F. Balabdaoui, A. E. Raftery, Probabilistic Forecasts, Calibration and Sharpness, *Journal of the Royal Statistical Society Series B: Statistical Methodology* 69 (2) (2007) 243–268. doi:[10.1111/j.1467-9868.2007.00587.x](https://doi.org/10.1111/j.1467-9868.2007.00587.x).
- [51] T. Gneiting, A. E. Raftery, Strictly Proper Scoring Rules, Prediction, and Estimation, *Journal of the American Statistical Association* 102 (477) (2007) 359–378. doi:[10.1198/016214506000001437](https://doi.org/10.1198/016214506000001437).
- [52] F. Tao, H. Zhang, A. Liu, A. Y. C. Nee, Digital Twin in Industry: State-of-the-Art, *IEEE Transactions on Industrial Informatics* 15 (4) (2019) 2405–2415. doi:[10.1109/tii.2018.2873186](https://doi.org/10.1109/tii.2018.2873186).

- [53] K. Fu, U. Ganslmeier, A. Mohapatra, R. Song, T. Hamacher, Real-world self-adaptive digital twin: An example of commercial air-source heat pump, *Energy and AI* 24 (2026) 100719. doi:[10.1016/j.egyai.2026.100719](https://doi.org/10.1016/j.egyai.2026.100719).
- [54] S. Kapoor, A. Narayanan, Leakage and the reproducibility crisis in machine-learning-based science, *Patterns* 4 (9) (2023) 100804. doi:[10.1016/j.patter.2023.100804](https://doi.org/10.1016/j.patter.2023.100804).
- [55] C. Rudin, Stop explaining black box machine learning models for high stakes decisions and use interpretable models instead, *Nature Machine Intelligence* 1 (5) (2019) 206–215. doi:[10.1038/s42256-019-0048-x](https://doi.org/10.1038/s42256-019-0048-x).
- [56] K. Aas, M. Jullum, A. Løland, Explaining individual predictions when features are dependent: More accurate approximations to Shapley values, *Artificial Intelligence* 298 (2021) 103502. doi:[10.1016/j.artint.2021.103502](https://doi.org/10.1016/j.artint.2021.103502).
- [57] A. Fisher, C. Rudin, F. Dominici, [All Models are Wrong, but Many are Useful: Learning a Variable’s Importance by Studying an Entire Class of Prediction Models Simultaneously](#), *Journal of Machine Learning Research* 20 (177) (2019) 1–81.
URL <https://jmlr.org/papers/v20/18-760.html>
- [58] Nixtla, [Temporal Fusion Transformer: NeuralForecast model documentation](#), implementation used: NeuralForecast 3.2.2; accessed 22 September 2026 (2026).
URL <https://nixtlaverse.nixtla.io/neuralforecast/models.tft.html>
- [59] Nixtla, [N-HiTS: NeuralForecast model documentation](#), implementation used: NeuralForecast 3.2.2; accessed 22 September 2026 (2026).
URL <https://nixtlaverse.nixtla.io/neuralforecast/models.nhits.html>
- [60] Google Research, [TimesFM 2.5 200M: model checkpoint and model card](#), model revision 1d952420fba87f3c6dee4f240de0f1a0fbc790e3; accessed 21 September 2026 (2026).
URL <https://huggingface.co/google/timesfm-2.5-200m-pytorch>
- [61] Google Research, [TimesFM: official implementation and covariate regression interface](#), timesFM 2.5 model class; package 3.0.2; accessed 22 September 2026 (2026).
URL <https://github.com/google-research/timesfm>

- [62] A. F. Ansari, O. Shchur, J. Küken, et al., [Chronos-2: From Univariate to Universal Forecasting](#) (2025). [arXiv:2510.15821](#).
URL <https://arxiv.org/abs/2510.15821>
- [63] Amazon Science, [Chronos forecasting: official implementation and Chronos-2 covariate interface](#), package chronos-forecasting 2.3.2; accessed 22 September 2026 (2026).
URL <https://github.com/amazon-science/chronos-forecasting>
- [64] T. Aksu, G. Woo, J. Liu, X. Liu, C. Liu, S. Savarese, C. Xiong, D. Sahoo, [GIFT-Eval: A Benchmark For General Time Series Forecasting Model Evaluation](#) (2024). [arXiv:2410.10393](#).
URL <https://arxiv.org/abs/2410.10393>
- [65] A. Zeng, M. Chen, L. Zhang, Q. Xu, Are Transformers Effective for Time Series Forecasting?, Proceedings of the AAAI Conference on Artificial Intelligence 37 (9) (2023) 11121–11128. [doi:10.1609/aaai.v37i9.26317](#).
- [66] C. Xu, Y. Xie, [Conformal prediction interval for dynamic time-series](#), in: Proceedings of the 38th International Conference on Machine Learning, Vol. 139 of Proceedings of Machine Learning Research, 2021, pp. 11559–11569.
URL <https://proceedings.mlr.press/v139/xu21h.html>
- [67] M. Zaffran, O. Feron, Y. Goude, J. Josse, A. Dieuleveut, [Adaptive Conformal Predictions for Time Series](#), in: Proceedings of the 39th International Conference on Machine Learning, Vol. 162 of Proceedings of Machine Learning Research, 2022, pp. 25834–25866.
URL <https://proceedings.mlr.press/v162/zaffran22a.html>
- [68] I. Gibbs, E. J. Candès, [Conformal Inference for Online Prediction with Arbitrary Distribution Shifts](#), Journal of Machine Learning Research 25 (162) (2024) 1–36.
URL <https://jmlr.org/papers/v25/22-1218.html>
- [69] R. J. Hyndman, A. B. Koehler, Another look at measures of forecast accuracy, International Journal of Forecasting 22 (4) (2006) 679–688. [doi:10.1016/j.ijforecast.2006.03.001](#).